

# Kafka Full Guide - All Details Included

## 1. What is Kafka?

Kafka is an open-source distributed event streaming platform designed to handle high-throughput, fault-tolerant, real-time data streams. It is widely used for event-driven architectures, data pipelines, and microservices communication.

## 2. Kafka Architecture & Components

Kafka's architecture consists of:

- **Producers:** Applications that send messages to Kafka.
- **Topics & Partitions:** Kafka topics are divided into partitions for scalability.
- **Consumers & Consumer Groups:** Consumers read messages in parallel.
- **Brokers & Clusters:** A Kafka cluster is composed of multiple brokers.
- **ZooKeeper:** Manages metadata, leader election, and broker coordination.

## 3. Kafka Producer

Kafka producers send messages to topics. They can control partitioning, acks, and message compression.

Example of sending messages using Kafka Producer CLI:

```
bin/kafka-console-producer.sh --topic my-topic --bootstrap-server localhost:9092
>Hello Kafka!
>Hello Again!
```

## 4. Kafka Consumer

Kafka consumers read messages from topics. They use consumer groups to parallelize message processing.

Example of consuming messages from the beginning:

```
bin/kafka-console-consumer.sh --topic my-topic --from-beginning --bootstrap-server localhost:9092
```

## 5. Creating a Kafka Topic

Kafka topics store messages and are divided into partitions. To create a topic:

```
bin/kafka-topics.sh --create --topic my-topic --bootstrap-server localhost:9092
--partitions 3 --replication-factor 1
```

## 6. Listing Kafka Topics

To list all available topics in Kafka:

```
bin/kafka-topics.sh --list --bootstrap-server localhost:9092
```

## 7. Kafka Streams API

Kafka Streams API is used for real-time stream processing. Example of transforming a stream:

```
java
KStream<String, String> source = builder.stream("input-topic");
KStream<String, String> transformed = source.mapValues(value -> value.toUpperCase());
transformed.to("output-topic");
```

## 8. Kafka Connect

Kafka Connect allows integration with external systems like databases, cloud storage, and more. Example of running a connector:

```
json
{
  "name": "jdbc-source-connector",
  "config": {
    "connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector",
    "connection.url": "jdbc:mysql://localhost:3306/mydb",
    "table.whitelist": "users",
    "mode": "incrementing",
    "incrementing.column.name": "id",
    "topic.prefix": "mysql-"
  }
}
```

## 9. Kafka Security

Kafka security includes authentication (SSL/SASL), authorization (ACLs), and encryption. Example: Enabling SSL encryption in Kafka:

```
listeners=SSL://:9093
ssl.keystore.location=/etc/kafka/secrets/kafka.keystore.jks
ssl.keystore.password=secret
```

## 10. Monitoring Kafka

Kafka monitoring can be done using JMX, Prometheus, and Grafana. Example: Exposing Kafka metrics for Prometheus:

```
KAFKA_OPTS="-javaagent:/usr/prometheus/jmx_prometheus_javaagent-0.3.1.jar=7071:/usr/prometheus/kafka.yml"
```

## 11. Kafka vs RabbitMQ vs ActiveMQ

Comparison of Kafka with RabbitMQ and ActiveMQ:

- **Kafka:** Best for event-driven, high-throughput messaging.
- **RabbitMQ:** Best for request-response messaging and task queues.
- **ActiveMQ:** Best for legacy enterprise messaging systems.

## 12. Running Kafka in Docker

Example docker-compose file to run Kafka and ZooKeeper in Docker:

```
yaml
version: '2'
services:
  zookeeper:
    image: wurstmeister/zookeeper
    ports:
      - "2181:2181"
  kafka:
    image: wurstmeister/kafka
    ports:
      - "9092:9092"
    environment:
      KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://localhost:9092
```

## 13. Kafka Message Delivery Semantics

Kafka provides three message delivery guarantees:

- **At most once:** Messages may be lost.
- **At least once:** Messages may be duplicated.
- **Exactly once:** Ensures no message loss and no duplication.

## 14. Advanced Kafka Optimization

Performance tuning for Kafka includes:

- **Producer:** Use batch size, linger.ms, and compression.
- **Consumer:** Increase fetch size and optimize offset commits.
- **Brokers:** Tune segment sizes and retention policies.

## 15. What's Next?

Once comfortable with Kafka basics, explore:

- Kafka Schema Registry (Avro, Protobuf, JSON)
- Kafka Transactions & Exactly-Once Processing
- Kafka Multi-Cluster Replication (MirrorMaker)
- Kafka in Kubernetes using Strimzi
- Kafka and Flink for real-time analytics